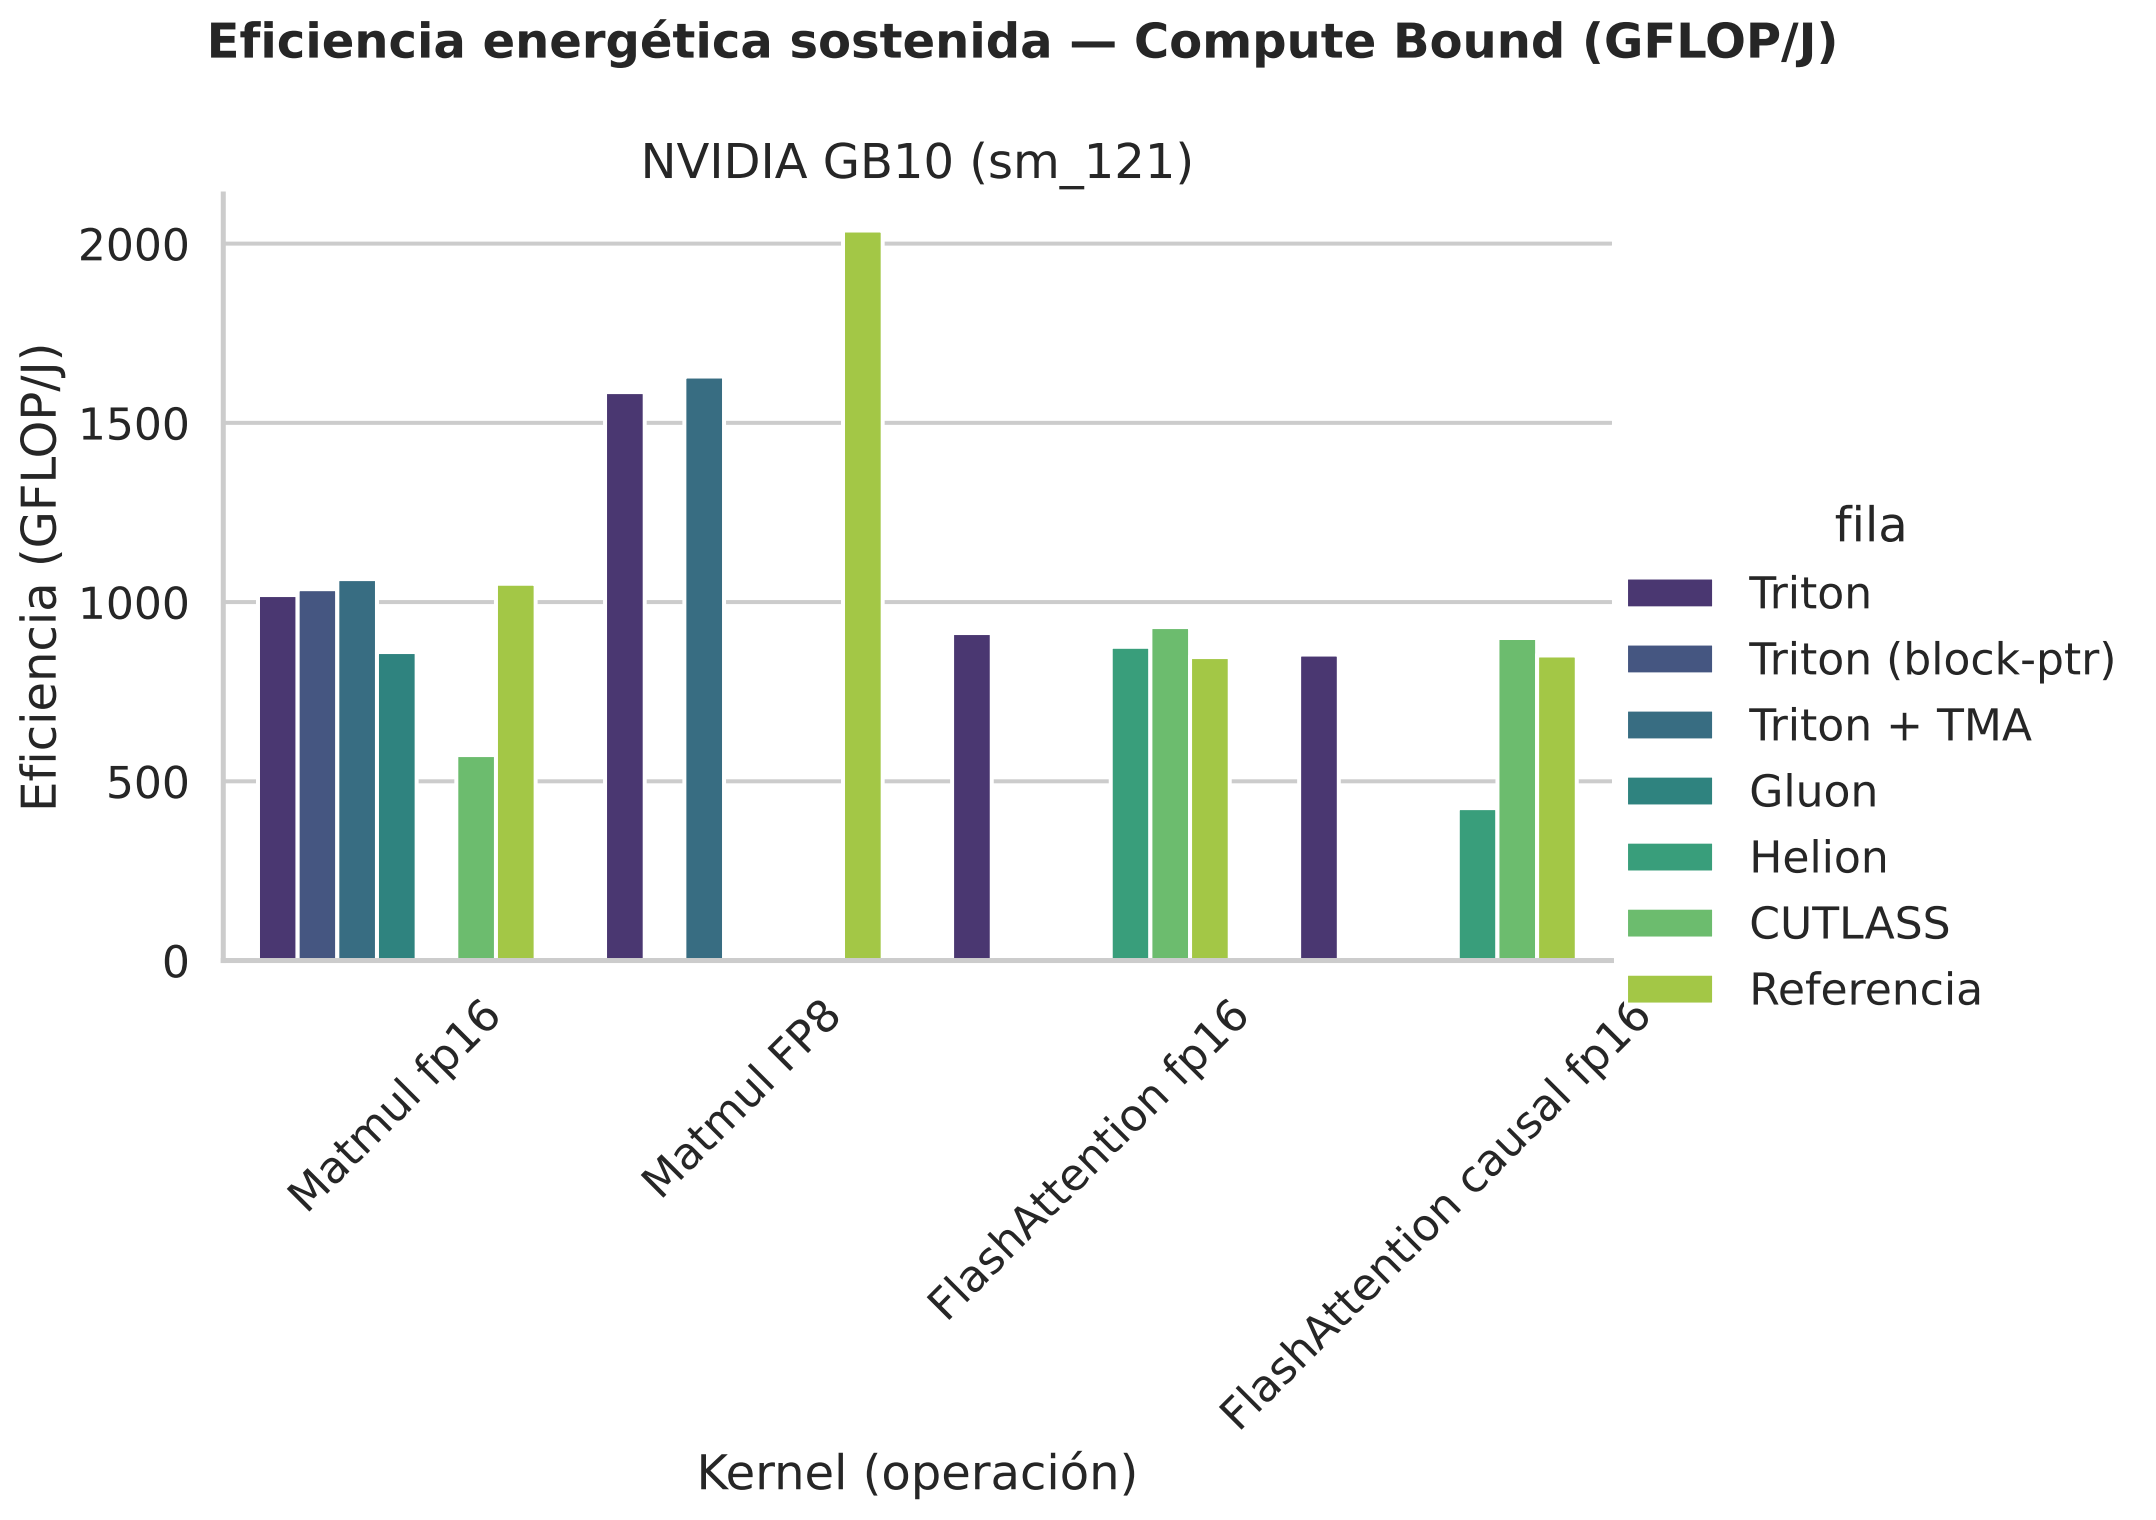
Eficiencia energética sostenida — Compute Bound (GFLOP/J)
NVIDIA GB10 (sm_121)
Eficiencia (GFLOP/J)
2000
1500
1000
500
0
Matmul fp16
Matmul FP8
FlashAttention fp16
FlashAttention causal fp16
Kernel (operación)
fila
Triton
Triton (block-ptr)
Triton + TMA
Gluon
Helion
CUTLASS
Referencia