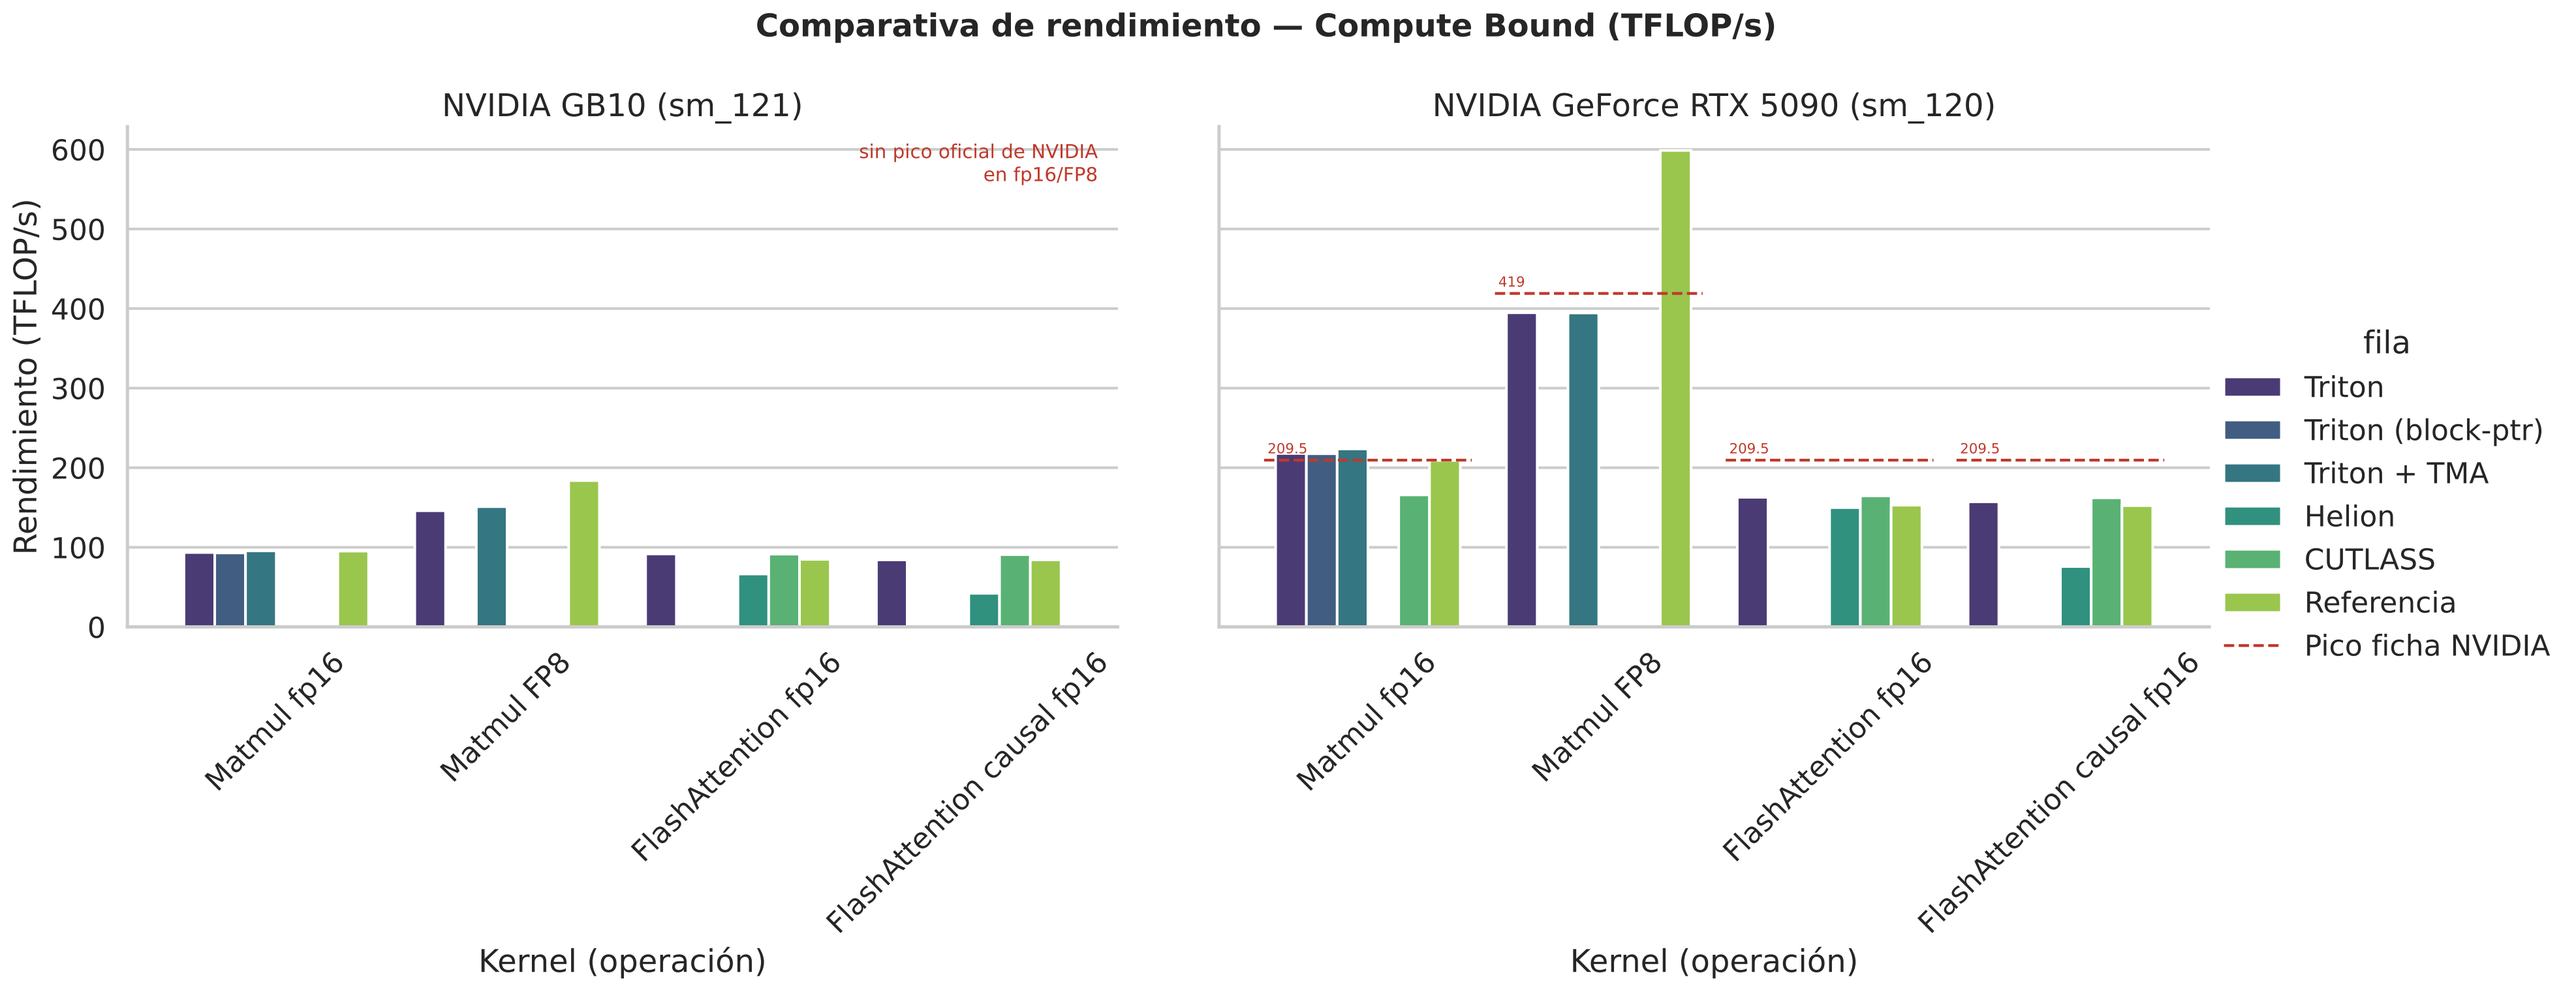
Comparativa de rendimiento — Compute Bound (TFLOP/s)
NVIDIA GB10 (sm_121)
NVIDIA GeForce RTX 5090 (sm_120)
Rendimiento (TFLOP/s)
0
100
200
300
400
500
600
sin pico oficial de NVIDIA
en fp16/FP8
209.5
419
209.5
209.5
Matmul fp16
Matmul FP8
FlashAttention fp16
FlashAttention causal fp16
Kernel (operación)
fila
Triton
Triton (block-ptr)
Triton + TMA
Helion
CUTLASS
Referencia
Pico ficha NVIDIA